

Homework-5 Report

Rohith Chandra Kandambeth

1. Methodology

a. Obtaining Queries

i. Description of the assigned queries

West African Ebola Epidemic: This epidemic refers to the largest outbreak of Ebola virus disease in history, primarily affecting several countries in West Africa from 2014 to 2016, most notably Liberia, Guinea, and Sierra Leone. Characterized by the transmission of the Ebola virus through direct contact with infected bodily fluids, the epidemic led to significant loss of life and economic disruption in the affected regions, prompting international public health responses.

H1N1 Swine Flu Pandemic: The H1N1 pandemic, which occurred in 2009, was caused by a novel influenza A virus. Originating in Mexico in early 2009 before spreading globally, this pandemic was characterized by the rapid transmission of H1N1, which contained a unique combination of influenza genes not previously identified in animals or people. Despite concerns about its severity, the pandemic's impact was mitigated through widespread vaccination efforts and public health measures, leading to fewer deaths than initially feared.

COVID-19: COVID-19, caused by the SARS-CoV-2 virus, was first identified in December 2019 in Wuhan, China, and quickly evolved into a global pandemic. Marked by symptoms ranging from mild respiratory issues to severe acute respiratory syndrome, COVID-19's highly contagious nature led to widespread health crises, economic disruptions, and significant mortality worldwide. The pandemic prompted global efforts in vaccination, public health measures, and changes to daily life to control its spread and impact.

ii. Explanation of the relevance to your vertical search engine topic

The relevance is high since the three queries above was used to create the index we are searching using the vertical search engine.

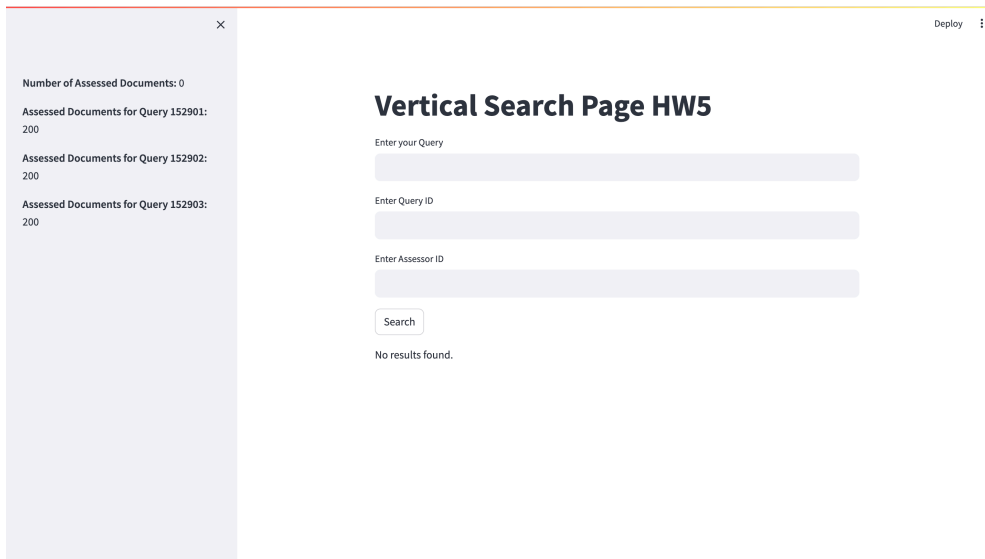
b. Assessment Graphical Interface

i. Features of the interface

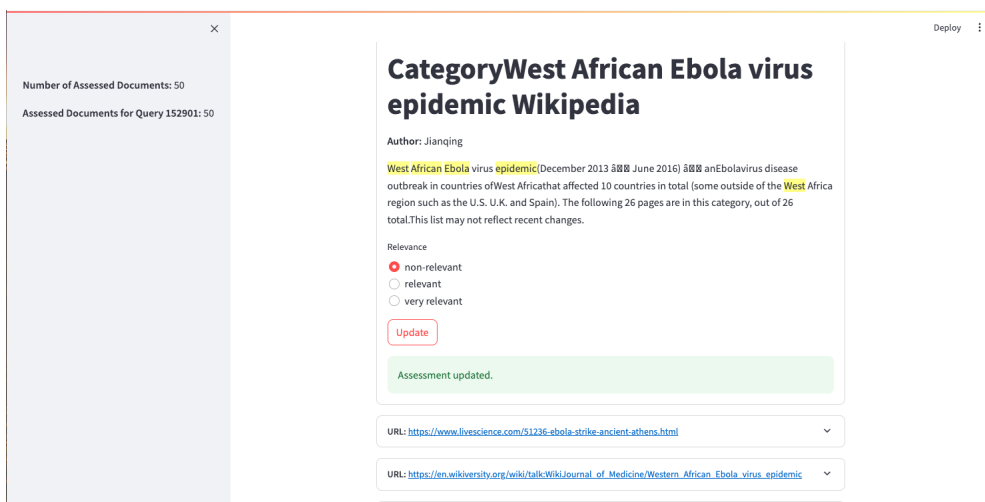
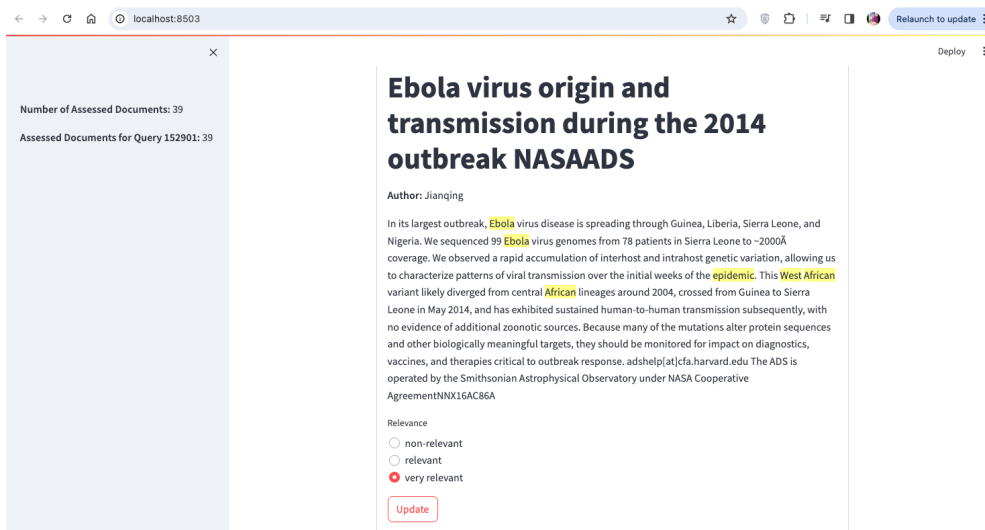
The interface allows the user to enter the Query, Query Id, and Accessor id. After the search button is clicked 250 results for the query is displayed using es search. Each result can be opened and marked as non-relevant, relevant, or very relevant. On the click of update button, the relevance of each document is saved to a file. The number of documents scored for each query is shown in the side of the interface.

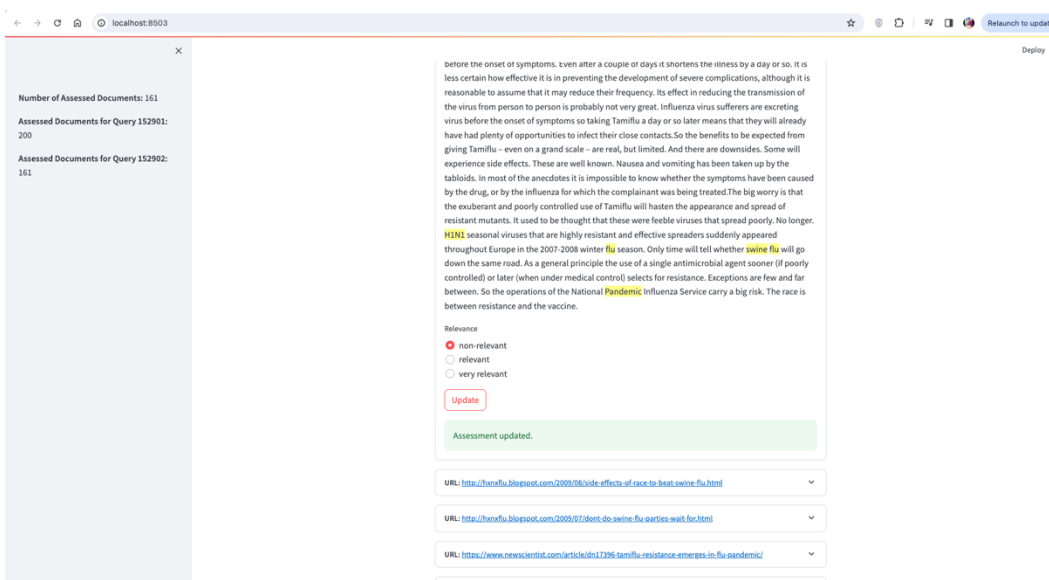
Additionally, the interface highlights the words in the query.

ii. Screenshots of the interface



Shows 200 for each query because I scored 600 documents already





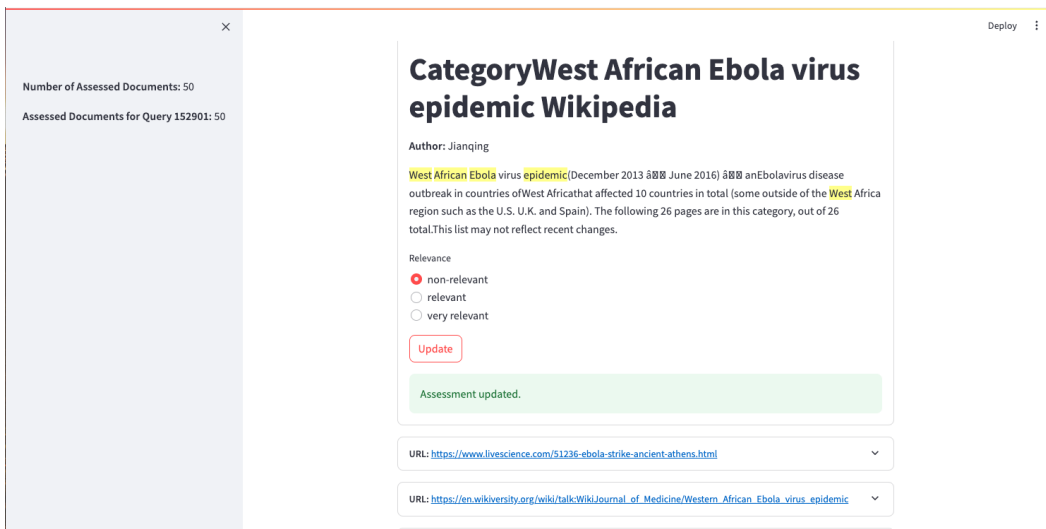
c. Manual Assessments

i. Approach to assessing the documents

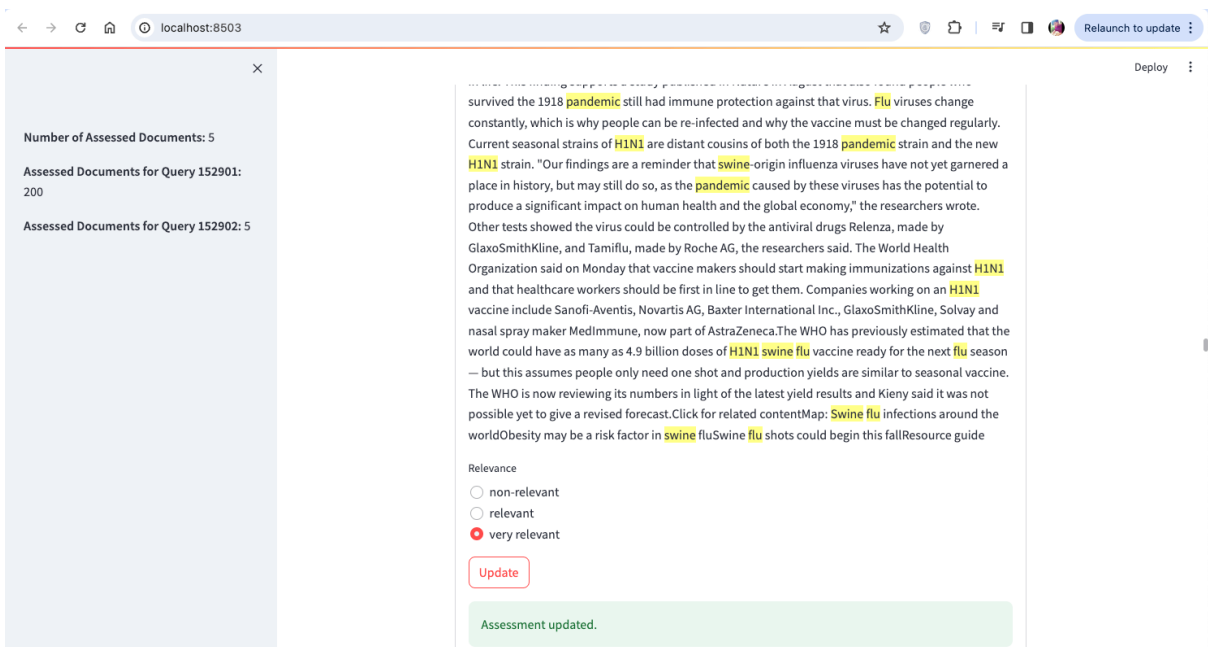
I assessed the documents based on the following factors

- Title of the document
- Number of words from the query that appears in the content (highlighted)
- Number of times the words (from the query) repeat (highlighted)
- If content is very small and does not impart any information, then it is rejected
- If the content gives out a lot of info, and the number of words repeat a lot, and title is related to the query then we can say the document is very relevant
- When the content is good enough, with some words from the query appearing couple of times then we can say the document is relevant.
- If none of the above complies then we can say the document is not relevant.

ii. Sample assessments with explanations



Here, the words from the query appears in content but it does not give out enough information therefore it is considered non-relevant



In the above case, there are lot of words from the second query that appear in the content, and it seems to repeat multiple times therefore we can consider it to be very relevant.

2. Implementation of trec_eval

a. Pseudocode

Inputs:

qrel_file: The file containing ground truth relevance judgments.

trec_file: The file containing the retrieval results from the system being evaluated.

print_all_queries: A boolean flag indicating whether to print metrics for each individual query.

Outputs:

Evaluation metrics such as MAP, nDCG, R-Precision, Precision at various cutoffs.
Steps:

1. Initialization
 - Set up a structure (dictionary) to store the ground truth relevance judgments (qrel).
 - Set up a structure (dictionary) to store the system's retrieval results (trec).
 - Define the metrics and cutoffs for evaluation.
2. Read the QREL File
 - For each line in the QREL file:
 - Parse the line to extract topic, doc_id, and relevance score.
 - Store the relevance score in the qrel structure, indexed by topic and doc_id.
3. Read the TREC File
 - For each line in the TREC file:
 - Parse the line to extract topic, doc_id, and score.
 - Store the score in the trec structure, indexed by topic and doc_id.
4. Evaluation Loop
 - For each topic in the trec structure:
 - If the topic is also in the qrel structure, proceed with evaluation; otherwise, skip the topic.
 - Sort the retrieved documents by score in descending order.
 - Calculate and accumulate the total counts for relevant documents, retrieved documents, and relevant retrieved documents.
5. Calculate Metrics at Cutoffs
 - For each predefined cutoff:
 - Calculate Precision, Recall, and F1 score for the top n documents up to the cutoff point.
6. Calculate R-Precision
 - Calculate the Precision after R (= number of relevant documents for a query) documents retrieved.
7. Calculate Average Precision
 - For each retrieved document, if it's relevant, accumulate the precision at the rank of this document.
 - Normalize the accumulated precision by the total number of relevant documents.

8. Calculate nDCG

- Calculate the ideal Discounted Cumulative Gain (iDCG) using the total number of relevant documents.
- Calculate the actual Discounted Cumulative Gain (DCG) using the relevance scores of the retrieved documents.
- Calculate nDCG by dividing DCG by iDCG.

9. Average Metrics Over All Topics

- Divide the accumulated metrics by the number of topics to get the average.

10. Printing Results

- If the `print_all_queries` flag is true, print the metrics for each query.
- Print the overall metrics averaged across all topics.

11. Main Function

- Parse command-line arguments to get `qrel_file`, `trec_file`, and `print_all_queries`.
- Call the `read_qrel` and `read_trec` functions to populate the data structures.
- Call the `evaluate` function to perform the evaluation and print the results.

b. Code snippets

```
3. def read_qrel(qrel_file):
4.     qrel = {}
5.     with open(qrel_file, 'r') as f:
6.         for line in f:
7.             topic, _, doc_id, rel = line.strip().split()
8.             if topic not in qrel:
9.                 qrel[topic] = {}
10.            qrel[topic][doc_id] = int(rel)
11.     return qrel
12.
13. def read_trec(trec_file):
14.     trec = {}
15.     with open(trec_file, 'r') as f:
16.         for line in f:
17.             topic, _, doc_id, _, score, _ = line.strip().split()
18.             if topic not in trec:
19.                 trec[topic] = {}
20.            trec[topic][doc_id] = float(score)
21.     return trec
22.
```

```
def evaluate(qrel, trec, print_all_queries):
    metrics = {'map': 0.0, 'ndcg': 0.0, 'r_prec': 0.0}
    cutoffs = [5, 10, 15, 20, 30, 100, 200, 500, 1000]
    for cutoff in cutoffs:
        metrics[f'p_{cutoff}'] = 0.0
        metrics[f'r_{cutoff}'] = 0.0
        metrics[f'f1_{cutoff}'] = 0.0

    num_topics = 0
    tot_num_ret = 0
```

```

tot_num_rel = 0
tot_num_rel_ret = 0

for topic in sorted(trec.keys()):
    if topic not in qrel:
        continue

    num_topics += 1
    retrieved_docs = sorted(trec[topic].items(), key=lambda x: x[1],
reverse=True)
    num_rel = sum(qrel[topic].values())
    num_ret = len(retrieved_docs)

    # Precision, Recall, F1 at cutoffs
    for cutoff in cutoffs:
        cutoff_docs = retrieved_docs[:cutoff]
        num_rel_ret = sum(qrel[topic].get(doc_id, 0) for doc_id, _ in
cutoff_docs)
        metrics[f'p_{cutoff}'] += num_rel_ret / cutoff
        metrics[f'r_{cutoff}'] += num_rel_ret / num_rel if num_rel > 0 else 0.0
        metrics[f'f1_{cutoff}'] += 2 * metrics[f'p_{cutoff}'] *
metrics[f'r_{cutoff}'] / (metrics[f'p_{cutoff}'] + metrics[f'r_{cutoff}']) if
metrics[f'p_{cutoff}'] + metrics[f'r_{cutoff}'] > 0 else 0.0

    # R-Precision
    r_prec_cutoff = num_rel
    r_prec_docs = retrieved_docs[:r_prec_cutoff]
    num_rel_ret = sum(qrel[topic].get(doc_id, 0) for doc_id, _ in r_prec_docs)
    metrics['r_prec'] += num_rel_ret / num_rel if num_rel > 0 else 0.0

    # Average Precision
    ap = 0.0
    num_rel_ret = 0
    for i, (doc_id, _) in enumerate(retrieved_docs):
        if qrel[topic].get(doc_id, 0) > 0:
            num_rel_ret += 1
            ap += num_rel_ret / (i + 1)
    ap /= num_rel if num_rel > 0 else 0.0
    metrics['map'] += ap

    # nDCG
    idcg = sum(1.0 / math.log2(i + 2) for i in range(num_rel))
    dcg = 0.0
    for i, (doc_id, _) in enumerate(retrieved_docs):
        rel = qrel[topic].get(doc_id, 0)
        if rel > 0:
            dcg += (2 ** rel - 1) / math.log2(i + 2)
    ndcg = dcg / idcg if idcg > 0 else 0.0
    metrics['ndcg'] += ndcg

    tot_num_ret += num_ret

```

```

    tot_num_rel += num_rel
    tot_num_rel_ret += num_rel_ret

    if print_all_queries:
        #print

# Average metrics over all queries
for metric in metrics:
    metrics[metric] /= num_topics

#print

```

3. Evaluation Results

a. Results of Manual Assessments

i. Compilation of assessments in QREL format

After scoring each document (600 documents), the data is stored in the below format in a file.

QueryID AssessorID DocID Grade

ii. Discussion on the assessment results

Most documents that are assessed are related to the query search. Some documents do not have any information or is misleading information containing the words from the query.

b. trec_eval Results

i. Presentation of results from your trec_eval

The results for trec eval for the hw3 crawled and searched data is shown below

```

Exact: 0.2303
(base) rohithram@Rohiths-MacBook-Air Results % Python trec_eval2.py qrels.txt query_result_es_builtin_hw3.txt

Queryid (Num): 3
Total number of documents over all queries
Retrieved: 3001
Relevant: 692
Rel_ret: 462
Interpolated Recall - Precision Averages:
at 0.00 2.0000
at 0.10 1.9667
at 0.20 1.8667
at 0.30 1.7778
at 0.40 1.5233
at 0.50 1.0467
at 0.60 0.4300
at 0.70 0.2210
at 0.80 0.2210
at 0.90 0.2210
at 1.00 0.2210
Average precision (non-interpolated) for all rel docs(averaged over queries)
0.6135
Precision:
At 5 docs: 2.0000
At 10 docs: 1.9667
At 15 docs: 1.9111
At 20 docs: 1.8667
At 30 docs: 1.7778
At 100 docs: 1.5233
At 200 docs: 1.0467
At 500 docs: 0.4300
At 1000 docs: 0.2210
R-Precision (precision after R (= num_rel for a query) docs retrieved):
Exact: 0.9140

```

ii. Comparison with provided trec-eval (if applicable)

Using the given trec eval file I got the below


```

Exact: 0.2303
(base) rohithram@Rohiths-MacBook-Air Results % perl trec_eval.pl qrels.adhoc.51-100.AP89.txt query_result_es_builtin2.txt
Error due to 23

Queryid (Num): 23
Total number of documents over all queries
Retrieved: 20079
Relevant: 1826
Rel_ret: 1060
Interpolated Recall - Precision Averages:
at 0.00 0.4763
at 0.10 0.3897
at 0.20 0.3144
at 0.30 0.2627
at 0.40 0.2192
at 0.50 0.1845
at 0.60 0.1400
at 0.70 0.1162
at 0.80 0.0892
at 0.90 0.0609
at 1.00 0.0226
Average precision (non-interpolated) for all rel docs(averaged over queries)
0.1925
Precision:
At 5 docs: 0.3217
At 10 docs: 0.3087
At 15 docs: 0.2957
At 20 docs: 0.2848
At 30 docs: 0.2725
At 100 docs: 0.1887
At 200 docs: 0.1341
At 500 docs: 0.0772
At 1000 docs: 0.0461
R-Precision (precision after R (= num_rel for a query) docs retrieved):
Exact: 0.2303

```

Using the trec eval code programmed by me I got the below

```

Exact: 0.2303
(base) rohithram@Rohiths-MacBook-Air Results % Python trec_eval2.py qrels.adhoc.51-100.AP89.txt query_result_es_builtin2.txt

Queryid (Num): 23
Total number of documents over all queries
Retrieved: 20079
Relevant: 1826
Rel_ret: 1060
Interpolated Recall - Precision Averages:
at 0.00 0.3217
at 0.10 0.3130
at 0.20 0.2848
at 0.30 0.2710
at 0.40 0.1887
at 0.50 0.1341
at 0.60 0.0772
at 0.70 0.0461
at 0.80 0.0461
at 0.90 0.0461
at 1.00 0.0461
Average precision (non-interpolated) for all rel docs(averaged over queries)
0.1924
Precision:
At 5 docs: 0.3217
At 10 docs: 0.3130
At 15 docs: 0.2928
At 20 docs: 0.2848
At 30 docs: 0.2710
At 100 docs: 0.1887
At 200 docs: 0.1341
At 500 docs: 0.0772
At 1000 docs: 0.0461
R-Precision (precision after R (= num_rel for a query) docs retrieved):
Exact: 0.2303

```

iii. Analysis of the results

The results produced by both the code is exactly the same telling us that the trec_eval code works perfectly.

```

Queryid (Num): 152901
R-Precision: 0.9951
Average Precision: 0.7173
nDCG: 1.4413
Precision@5: 2.0000
Recall@5: 0.0488
F1@5: 0.0952
Precision@10: 1.9000
Recall@10: 0.0927
F1@10: 0.1767
Precision@15: 1.8667
Recall@15: 0.1366
F1@15: 0.2545
Precision@20: 1.7000
Recall@20: 0.1659
F1@20: 0.3022
Precision@30: 1.6000
Recall@30: 0.2341
F1@30: 0.4085
Precision@100: 1.3200
Recall@100: 0.6439
F1@100: 0.8656
Precision@200: 1.0200
Recall@200: 0.9951
F1@200: 1.0074
Precision@500: 0.4100
Recall@500: 1.0000
F1@500: 0.5816
Precision@1000: 0.2050
Recall@1000: 1.0000
F1@1000: 0.3402

```

The system appears to have a strong ability to rank relevant documents highly, but there may be a significant number of relevant documents that are not being retrieved within the top results, based on the recall values. There also seems to be a potential issue with the calculation or interpretation of some metrics, specifically Precision@5 and nDCG, which should be reviewed for accuracy

c. Precision-Recall Curves (For MS Students)

i. Description of the method to create precision-recall plots

1. Inside the loop that iterates over each topic in the `trec` dictionary, two empty lists, `precision` and `recall`, are initialized to store the precision and recall values at different cutoff points.
2. For each cutoff value in the `cutoffs` list, the following steps are performed:
 - The retrieved documents up to the current cutoff are obtained using `retrieved_docs[:cutoff]` and stored in `cutoff_docs`.
 - The number of relevant documents retrieved up to the current cutoff is calculated using `sum(qrel[topic].get(doc_id, 0) for doc_id, _ in cutoff_docs)` and stored in `num_rel_ret`.
 - The precision at the current cutoff is calculated as `num_rel_ret / cutoff` and appended to the `precision` list.
 - The recall at the current cutoff is calculated as `num_rel_ret / num_rel` if `num_rel > 0`, otherwise it is set to 0.0. The recall value is appended to the `recall` list.
3. After calculating the precision and recall values for all cutoff points, the code proceeds to plot the precision-recall curve using the `matplotlib` library.
4. A new figure is created with a specified size using `plt.figure(figsize=(8, 6))`.
5. The precision-recall curve is plotted using `plt.plot(recall, precision, marker='o', linestyle='-', linewidth=2)`. This function takes the recall values as the x-coordinates and the precision values as the y-coordinates. The `marker` parameter is set to 'o' to

display markers at each data point, `linestyle` is set to '-' for a solid line, and linewidth` is set to 2 to specify the thickness of the line.`

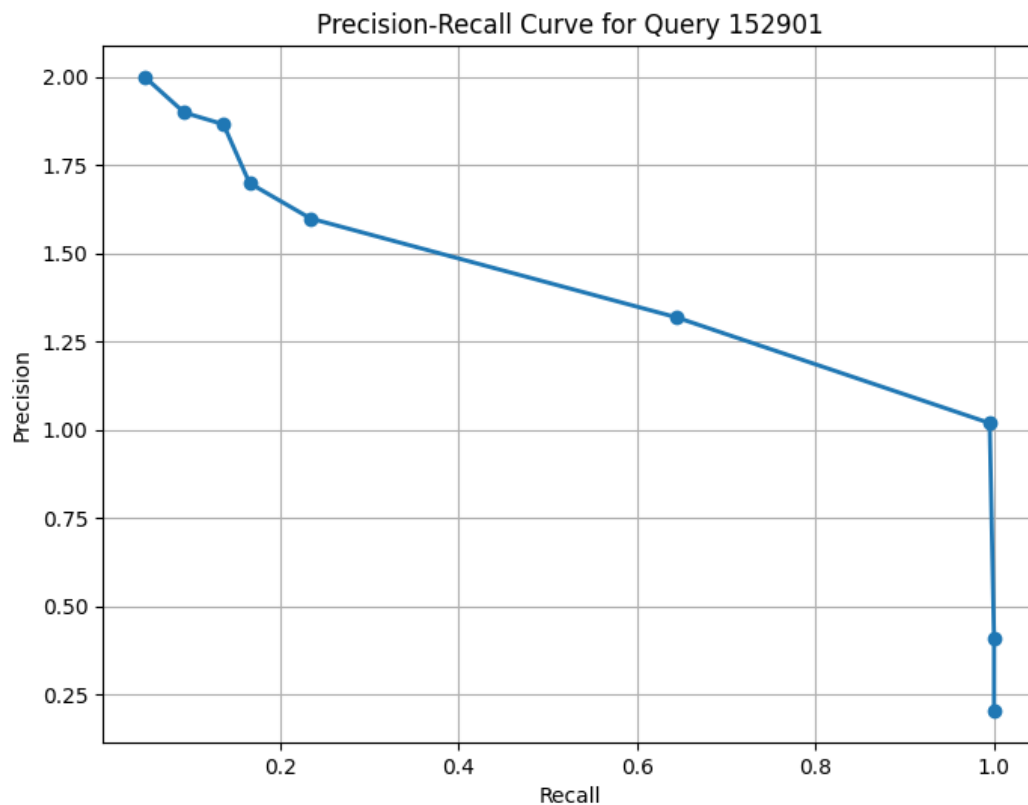
6. Labels for the x-axis and y-axis are added using `plt.xlabel('Recall')` and plt.ylabel('Precision')`, respectively.`

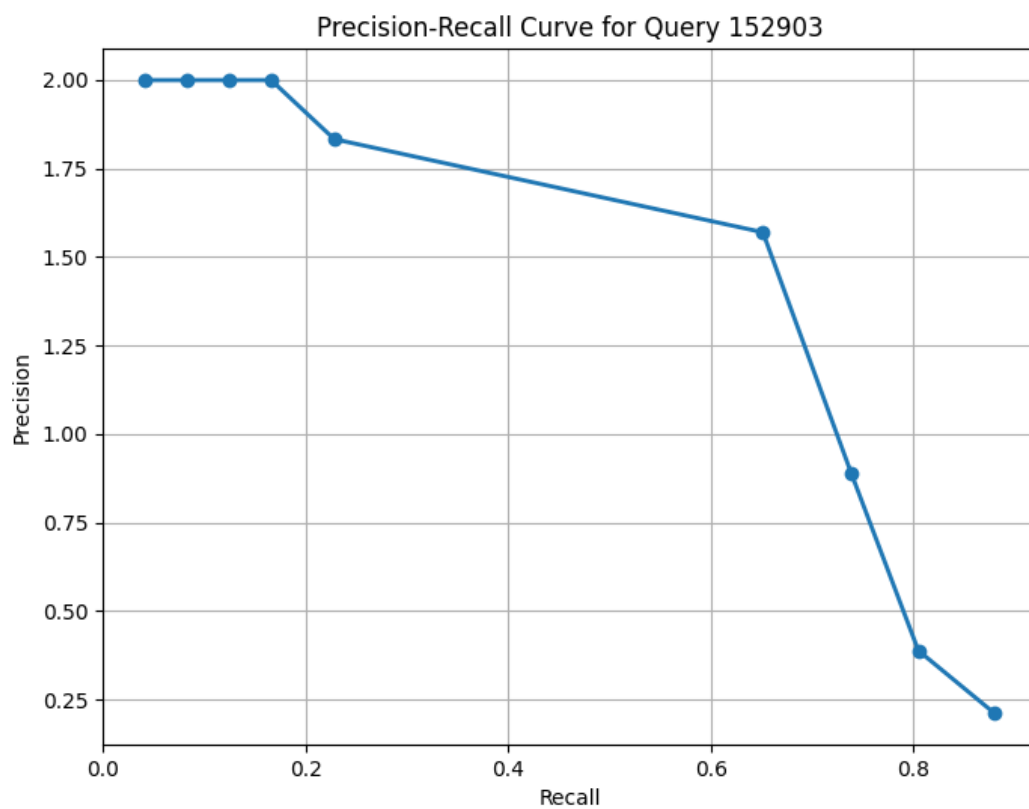
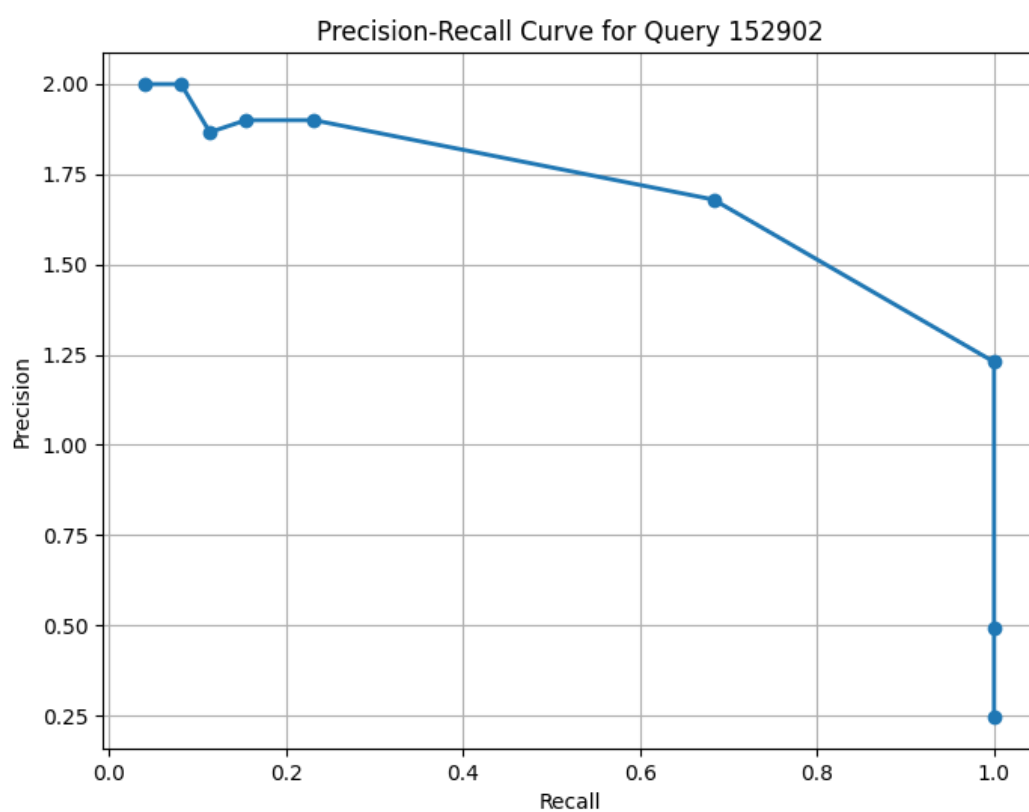
7. A title is added to the plot using `plt.title(f'Precision-Recall Curve for Query {topic}')`, which includes the query topic number.

8. A grid is added to the plot using `plt.grid(True)` to improve readability.`

9. Finally, the plot is saved as a PNG file using `plt.savefig(f'precision_recall_query_{topic}.png')`, where the filename includes the query topic number.

ii. Graphs for each query





Query 152901:

The precision-recall curve shows a steep drop in precision as recall increases. This indicates that the retrieval system returns many non-relevant documents early in the ranked list, resulting in lower precision at higher recall levels.

Query 152902:

The precision-recall curve exhibits a gradual decline in precision as recall increases. The system maintains relatively high precision at lower recall levels but precision decreases as more documents are retrieved. This suggests a better balance between precision and recall compared to Query 152901.

Query 152903:

The precision-recall curve demonstrates a rapid decline in precision at lower recall levels, followed by a more gradual decrease. The system achieves high precision for the top-ranked documents but struggles to maintain precision as recall increases. The curve shape indicates room for improvement in retrieving relevant documents at higher recall levels.

The precision-recall curves provide insights into the trade-off between precision and recall for each query. Query 152902 shows the best performance, maintaining higher precision at various recall levels compared to the other two queries.

6. Extra Credit (If Attempted)

- Description of the extra credit task(s) undertaken
Created my own web interface
Qualities – Simple, Fast and Efficient Webpage. Highlights the words from the query in the content making it easy to assess for relevance